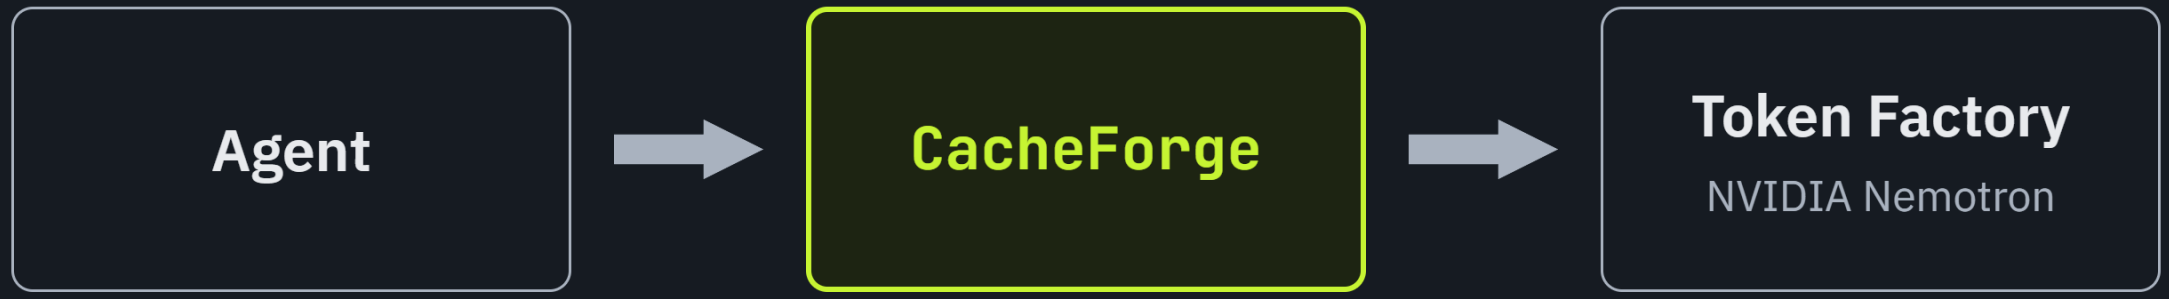
CacheForge:
a drop-in gateway

Run 2.2–4.6× more agent sessions per dollar on Nebius Token Factory.

Test logs are re-sent on every turn. Input is almost the whole bill.

13 of 15

sessions ran out of context without CacheForge (Nemotron 3 Nano, 32,768-token window)




```
Keeps the failing test, its assertion and setup lines, verbatim.
```

One test log: 40,141 → 9,662 characters.

2.2–4.6×

as many agent sessions per dollar, in all 7 comparisons

55–78%

lower cost per session

Time to first token

17–33%

lower

Ran out of context

0 of 105

with CacheForge, 3 of 90 without

Largest request

256,167 → 28,436

tokens

1 win	5 ties	1 loss

Head-to-head: no CacheForge 52, v2 60, v2.1 58 bugs fixed. No significant difference.

> **The claim is cost, not more bugs fixed.**

Found and fixed

- A fake first result (41 vs 7). The real result: 58 vs 37.
- Condenser v1 hid the clue: 37 → 13. v2 keeps the setup lines.
- v2 was slow on a local GPU: 229 s → 331 s. v2.1 fixes that.
- A pre-registered rule fired against my preferred version. I followed it.

Cheaper sessions mean more agents get run.

Smaller requests. Faster first token.

Feedback: about 90% of input was already cached, but cached input isn’t priced lower.

github.com/adamm285-dev/CacheForge

MIT license · Change one setting: your agent's base URL

Every result file and pre-registration is in the repo.